methylKit: User guide

Altuna Akalin

December 21, 2011

1 Introduction

In this manual, we will show how to use the methylKit package. methylKit is an R package for analysis and annotation of DNA methylation information obtained by high-throughput bisulfite sequencing. The package is designed to deal with sequencing data from RRBS and its variants. But it can potentially handle whole-genome bisulfite sequencing data if proper input format is provided.

2 Basics

2.1 Reading the methylation call files

We start by reading in the methylation call data from bisulfite sequencing with `read` function. Reading in the data this way will return a methylRawList object which stores methylation information per sample.

```
> library(methylKit)
> file.list=list( system.file("extdata", "test1.myCpG.txt", package = "methylKit"),
+                 system.file("extdata", "test2.myCpG.txt", package = "methylKit"),
+                 system.file("extdata", "control1.myCpG.txt", package = "methylKit"),
+                 system.file("extdata", "control2.myCpG.txt", package = "methylKit") )
> # read the files to a methylRawList object: myobj
> myobj=read(file.list,
+            sample.id=list("test1","test2","ctrl1","ctrl2"),
+            assembly="hg18",
+            pipeline="amp",
+            treatment=c(1,1,0,0))
>
>
```

2.2 Descriptive statistics on samples

Since we read the data now, we can check the basic stats about the methylation data such as coverage and percent methylation. We now have a `methylRawList` object which contains methylation information per sample. The following command prints out percent methylation statistics for second sample: "test2"

```
> getMethylationStats(myobj[[2]],plot=F,both.strands=F)
```

```
methylation statistics per base
summary:
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   0.00   20.00   82.79   63.17   94.74  100.00
percentiles:
        0%        10%        20%        30%        40%        50%        60%        70%
  0.00000    0.00000    0.00000   48.38710   70.00000   82.78556   90.00000   93.33333
       80%        90%        95%        99%      99.5%      99.9%       100%
 96.42857 100.00000 100.00000 100.00000 100.00000 100.00000 100.00000
```

The following command plots the histogram for percent methylation distribution.The figure below is the histogram and numbers on bars denote what percentage of locations are contained in that bin. Typically, percent methylation histogram should have two peaks on both ends. In any given cell, any given base are either methylated or not. Therefore, looking at many cells should yield a similar pattern where we see lots of locations with high methylation and lots of locations with low methylation.

```
> library("graphics")
> getMethylationStats(myobj[[2]],plot=T,both.strands=F)
```

Histogram of % methylation
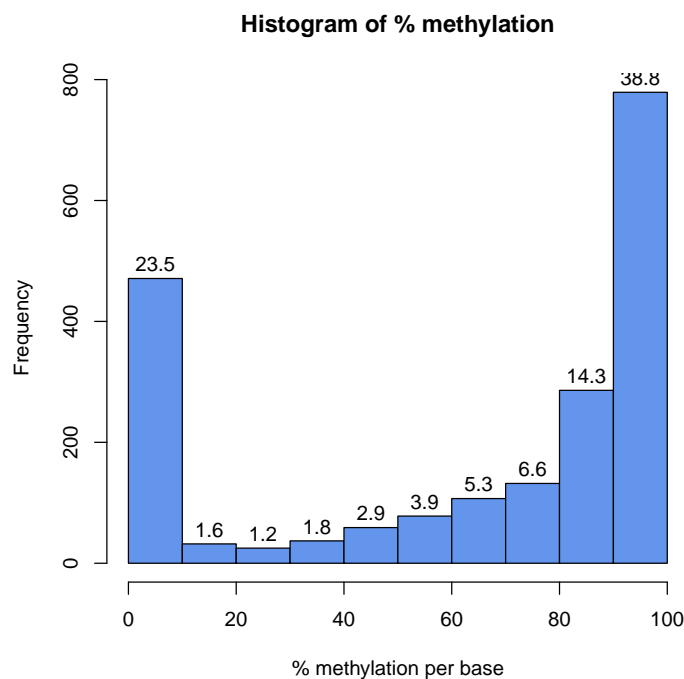
We can also plot the read coverage per base information in a similar way, again numbers on bars denote what percentage of locations are contained in that bin. Experiments that are highly suffering from PCR duplication bias will have a secondary peak towards the right hand side of the histogram.

```
> library ("graphics")
> getCoverageStats(myobj[[2]],plot=T,both.strands=F)
```

Histogram of CpG coverage
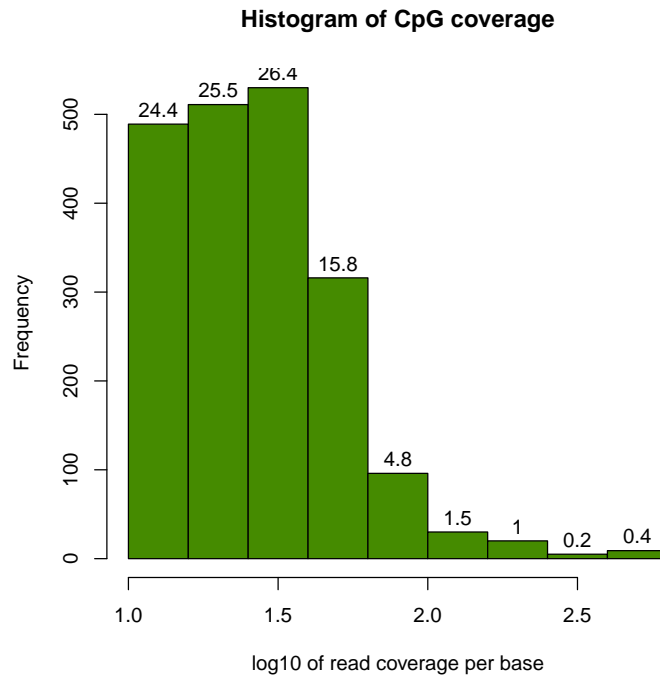
2.3 Filtering samples based on read coverage

It might be useful to filter samples based on coverage. Particularly, if our samples are suffering from PCR bias it would be useful to discard bases with very high read coverage. Furthermore, we would also like to discard bases that have low read coverage, a high enough read coverage will increase the power of the statistical tests. The code below filters a `methylRawList` and discards bases that have coverage below 10X and also discards the bases that have more than 99.9th percentile of coverage in each sample.

```
> filtered.myobj=filterByCoverage(myobj,lo.count=10,lo.perc=NULL,
+                                 hi.count=NULL,hi.perc=99.9)
```

3 Comparative analysis

3.1 Merging samples

In order to do further analysis, we will need to get the bases covered in all samples. The following function will merge all samples to one object for base-pair locations that are covered in all samples. Setting destrand=TRUE (the default is FALSE) will merge reads on both strands of a CpG dinucleotide. This provides better coverage, but only advised when looking at CpG methylation (for CpH methylation this will cause wrong results in subsequent analyses). This operation will return a methylBase object which will be our main object for all comparative analysis.

```
> methidh=unite(myobj,destrand=FALSE)
```

Let us take a look at the data content of methylBase object:

```
> head(methidh)
```

```
               id   chr     start       end strand coverage1 numCs1 numTs1
1 chr21.10011833 chr21 10011833 10011833      +       174    173      1
2 chr21.10011841 chr21 10011841 10011841      +       173    164      9
3 chr21.10011855 chr21 10011855 10011855      +       175    175      0
4 chr21.10011858 chr21 10011858 10011858      +       175    131     44
5 chr21.10011861 chr21 10011861 10011861      +       174    147     27
6 chr21.10011872 chr21 10011872 10011872      +       167    160      7
  coverage2 numCs2 numTs2 coverage3 numCs3 numTs3 coverage4 numCs4 numTs4
1        18     18      0        40     34      6        14     14      0
2        20     19      1        40     18     22        14      8      6
3        21     21      0        39     29     10        14     12      2
4        21     20      1        39     31      8        13      8      5
5        20     15      5        39     13     26        13      9      4
6        20     19      1        39     34      5        14      8      6
```
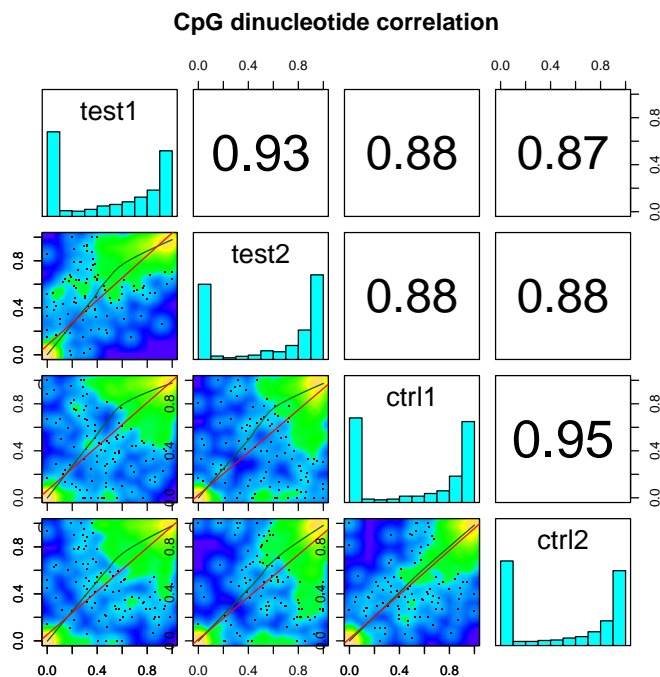
3.2 Sample Correlation

We can check the correlation between samples using `getCorrelation`. This function will either plot scatter plot and correlation coefficients or just print a correlation matrix

```
> getCorrelation(methidh,plot=T)
```

```
          test1      test2      ctrl1      ctrl2
test1 1.0000000 0.9252530 0.8767865 0.8737509
test2 0.9252530 1.0000000 0.8791864 0.8801669
ctrl1 0.8767865 0.8791864 1.0000000 0.9465369
ctrl2 0.8737509 0.8801669 0.9465369 1.0000000
```

CpG dinucleotide correlation

3.3 Clustering samples

We can cluster the samples based on the similarity of their methylation profiles.
The following function will cluster the samples and draw a dendrogram.

```
> clusterSamples(methidh, dist="correlation", method="ward", plot=TRUE)

Call:
hclust(d = d, method = HCLUST.METHODS[hclust.method])

Cluster method   : ward
Distance         : pearson
Number of objects: 4
```

CpG dinucleotide methylation clustering

Samples
Distance method: "correlation"; Clustering method: "ward"

Setting the plot=FALSE will return a dendrogram object which can be manipulated by users or fed in to other user functions that can work with dendrograms.

```
> hc = clusterSamples(methidh, dist="correlation", method="ward", plot=FALSE)
```

We can also do a PCA analysis on our samples. The following function will plot a scree plot for importance of components.

```
> PCASamples(methidh, screeplot=TRUE)
```

```
Importance of components:
                          Comp.1     Comp.2     Comp.3     Comp.4
Standard deviation      1.9211651 0.42545636 0.2735654 0.23081050
Proportion of Variance  0.9227188 0.04525328 0.0187095 0.01331837
Cumulative Proportion   0.9227188 0.96797213 0.9866816 1.00000000
```

CpG dinucleotide methylation PCA Screeplot
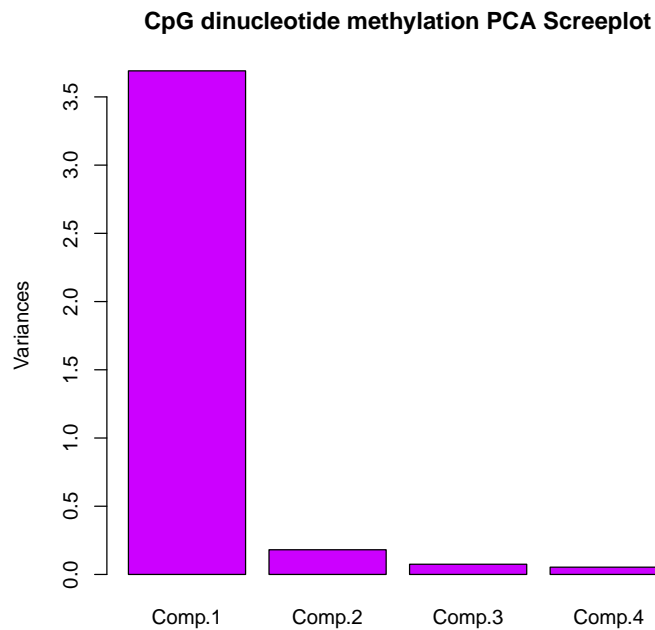
We can also plot PC1 and PC2 axis and a scatter plot of our samples on those axis which will reveal how they cluster.

```
> PCASamples(methidh)

Importance of components:
                           Comp.1     Comp.2    Comp.3     Comp.4
Standard deviation      1.9211651 0.42545636 0.2735654 0.23081050
Proportion of Variance  0.9227188 0.04525328 0.0187095 0.01331837
Cumulative Proportion   0.9227188 0.96797213 0.9866816 1.00000000
```

CpG dinucleotide methylation PCA Analysis

3.4 Tiling windows analysis

For some situations, it might be desirable to summarize methylation information over tiling windows rather than doing base-pair resolution analysis. `methylKit` provides functionality to do such analysis. The function below tiles the genome with windows 1000bp length and 1000bp step-size and summarizes the methylation information on those tiles. In this case, it returns a `methylRawList` object which can be fed into `unite` and `calculateDiffMeth` functions to get differentially methylated regions

```
> tiles=tileMethylCounts(myobj,win.size=1000,step.size=1000)
> head(tiles[[1]])
```

```
                      id   chr   start      end strand coverage numCs numTs
1 chr21.9764001.9765000 chr21 9764001 9765000      *        24     3    21
2 chr21.9820001.9821000 chr21 9820001 9821000      *        13     0    13
3 chr21.9837001.9838000 chr21 9837001 9838000      *        11     0    11
4 chr21.9849001.9850000 chr21 9849001 9850000      *       124    90    34
5 chr21.9853001.9854000 chr21 9853001 9854000      *        34    22    12
6 chr21.9860001.9861000 chr21 9860001 9861000      *        39    38     1
```

3.5 Finding differentially methylated bases or regions

`calculateDiffMeth()` function is the main function to calculate differential methylation. Depending on the sample size per each set it will either use Fisher's exact or logistic regression to calculate P-values. P-values will be adjusted to Q-values using SLIM method[1].

```
> myDiff=calculateDiffMeth(methidh)
```

After q-value calculation, we can select the differentially methylated regions/bases based on q-value and percent methylation difference cutoffs. Following bit selects the bases that have q-value<0.01 and percent methylation difference larger than 25%. If you specify type="hyper" or type="hypo" options, you will get hyper-methylated or hypo-methylated regions/bases.

```
> # get hyper methylated bases
> myDiff25p.hyper=get.methylDiff(myDiff,difference=25,qvalue=0.01,type="hyper")
> #
> #
> # get all differentially methylated bases
> myDiff25p=get.methylDiff(myDiff,difference=25,qvalue=0.01,type="all")
```

We can also visualize the distribution of hypo/hyper-methylated bases/regions per chromosome using the following function. In this case, the example set includes only one chromosome. The list shows percentages of hypo/hyper methylated bases over all the covered bases in a given chromosome.

```
> diffMethPerChr(myDiff,plot=FALSE,qvalue.cutoff=0.01, meth.cutoff=25)

$diffMeth.per.chr
    chr number.of.hypomethylated percentage.of.hypomethylated
1 chr21                       59                     6.126687
  number.of.hypermethylated percentage.of.hypermethylated
1                        75                     7.788162


$diffMeth.all
  percentage.of.hypermethylated number.of.hypermethylated
1                     7.788162                        75
  percentage.of.hypomethylated number.of.hypomethylated
1                     6.126687                        59
```

4 Annotating differentially methylated bases or regions

We can annotate our differentially methylated regions/bases based on gene annotation. In this example, we read the gene annotation from a bed file and annotate our differentially methylated regions with that information. This will tell us what percentage of our differentially methylated regions are on promoters/introns/exons/intergenic region. Similar gene annotation can be fetched using GenomicFeatures package available from Bioconductor.org.

```
> gene.obj=read.transcript.features(system.file("extdata", "refseq.hg18.bed.txt",
+                                          package = "methylKit"))
> #
> # annotate differentially methylated Cs with promoter/exon/intron using annotation data
> #
> annotate.WithGenicParts(myDiff25p,gene.obj)
```

```
summary of target set annotation with genic parts
133 rows in target set
--------------
--------------
percentage of target features overlapping with annotation :
  promoter       exon      intron intergenic
  27.81955    15.03759    34.58647   57.14286


percentage of target features overlapping with annotation (with promoter>exon>intron prece
  promoter       exon      intron intergenic
  27.81955     0.00000    15.03759   57.14286


percentage of annotation boundaries with feature overlap :
   promoter         exon        intron
0.018129079 0.001589593 0.010038738


summary of distances to the nearest TSS :
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      5     828   45160   52030   94640  313500
```

Similarly, we can read the CpG island annotation and annotate our differentially methylated bases/regions with them.

```
> # read the shores and flanking regions and name the flanks as shores
> # and CpG islands as CpGi
> cpg.obj=read.feature.flank(system.file("extdata", "cpgi.hg18.bed.txt",
+                                         package = "methylKit"),
+                     feature.flank.name=c("CpGi","shores"))
> #
> #
> diffCpGann=annotate.WithFeature.Flank(myDiff25p,cpg.obj$CpGi,cpg.obj$shores,
+                                       feature.name="CpGi",flank.name="shores")
```

4.1   Regional analysis

We can also summarize methylation information over a set of defined regions such as promoters or CpG islands. The function below summarizes the methylation information over a given set of promoter regions and outputs a `methylRaw` or `methylRawList` object depending on the input.

```
> promoters=regionCounts(myobj,gene.obj$promoters)
> head(promoters[[1]])

                         id   chr    start      end strand coverage numCs
1 chr21.17806094.17808094.NA chr21 17806094 17808094      +     1834     7
2 chr21.10119796.10121796.NA chr21 10119796 10121796      -       79    44
3 chr21.10011791.10013791.NA chr21 10011791 10013791      -     3697  2982
4 chr21.10119808.10121808.NA chr21 10119808 10121808      -       79    44
```

```
5 chr21.15357997.15359997.NA chr21 15357997 15359997          -      8613    16
6 chr21.16023366.16025366.NA chr21 16023366 16025366          +      6296     5
  numTs
1  1827
2    35
3   715
4    35
5  8594
6  6291
```

4.2   Convenience functions for annotation objects

After getting the annotation of differentially methylated regions, we can get the distance to TSS and nearest gene name using the `getAssociationWithTSS` function.

```
> diffAnn=annotate.WithGenicParts(myDiff25p,gene.obj)
> # target.row is the row number in myDiff25p
> head(getAssociationWithTSS(diffAnn))

     target.row dist.to.feature feature.name feature.strand
60            1             951    NM_199260              -
60.1          2             931    NM_199260              -
60.2          3             838    NM_199260              -
60.3          4             828    NM_199260              -
60.4          5             802    NM_199260              -
60.5          6             723    NM_199260              -
```

It is also desirable to get percentage/number of differentially methylated regions that overlap with intron/exon/promoters

```
> getTargetAnnotationStats(diffAnn,percentage=TRUE,precedence=TRUE)

   promoter        exon       intron intergenic
   27.81955     0.00000     15.03759   57.14286
```

We can also plot the percentage of differentially methylated bases overlapping with exon/intron/promoters

```
> plotTargetAnnotation(diffAnn,precedence=TRUE,
+     main="differential methylation annotation")
```

differential methylation annotation

We can also plot the CpG island annotation the same way. The plot below shows what percentage of differentially methylated bases are on CpG islands, CpG island shores and other regions.

```
> plotTargetAnnotation(diffCpGann,col=c("green","gray","white"),
+         main="differential methylation annotation")
```

differential methylation annotation

It might be also useful to get percentage of intron/exon/promoters that overlap with differentially methylated bases.

```
> getFeatsWithTargetsStats(diffAnn,percentage=FALSE)

promoter    exon   intron
       5       6       34
```

References

[1] Hong-Qiang Wang, Lindsey K Tuominen, and Chung-Jui Tsai. SLIM: a sliding linear model for estimating the proportion of true null hypotheses in datasets with dependence structures. *Bioinformatics (Oxford, England)*, 27(2):225–31, January 2011.